


COGS 108 Week 6

A04/A07

Nov 6, 2023





AGENDA FOR TODAY



LOGISTICS



SENTIMENT
ANALYSIS



DISCUSSION
LAB 5



LOGISTICS




DUE DATES

- A2 is due Nov 08, 11:59PM (this Wednesday)
- D5 is due Monday, Nov 13, 11:59PM | Extra Days!!
- Project CheckPoint #1 coming up on Nov 15, 11:59PM
 - Understand the feedback received in the project proposal.
 - Use TA/Professor OH to discuss on the feedback and next steps.



Sentiment Analysis





Sentiment Analysis



Sentiment analysis, often referred to as opinion mining, is a field of Natural Language Processing (NLP) that aims to extract and interpret emotions and opinions from textual data. By analyzing words, phrases, or even entire documents, sentiment analysis can determine whether the expressed sentiment is positive, negative, neutral, or sometimes even more nuanced emotions like joy, anger, or surprise.





Several Ways to Perform Sentiment Analysis

Rule-based:

The rule-based method uses predefined word lists (lexicons) to detect and score emotions in text. Marketers give scores to these words based on their emotional impact. To gauge a sentence's sentiment, the software tallies these scores. The total score then decides if the sentiment is positive, negative, or neutral.



Machine Learning:

Using ML techniques like neural networks, this method teaches software to detect emotions in text. It involves training a sentiment model on known data to accurately predict sentiments in unfamiliar data.





DISCUSSION LAB 5: INFERENCE





Part I : Data & Wrangling

The `apply()` method allows you to apply a function along one of the axis of the DataFrame, default 0, which is the index (row) axis.

Syntax

```
dataframe.apply(func, axis, raw, result_type, args, kwds)
```

Parameters

The `axis`, `raw`, `result_type`, and `args` parameters are keyword arguments.

Parameter	Value	Description
<i>func</i>		Required. A function to apply to the DataFrame.
axis	0 1 'index' 'columns'	Optional, Which axis to apply the function to. default 0.
raw	True False	Optional, default False. Set to true if the row/column should be passed as an ndarray object
result_type	'expand' 'reduce' 'broadcast' None	Optional, default None. Specifies how the result will be returned
args	<i>a tuple</i>	Optional, arguments to send into the function
kwds	<i>keyword arguments</i>	Optional, keyword arguments to send into the function



Part II : EDA

```
fig =  
pd.plotting.scatter_matrix(  
df[['column_1',  
column_2]])
```

A scatter matrix is useful because it allows you to visualize the relationship between multiple variables in a dataset at once.

pandas.plotting.scatter_matrix

```
pandas.plotting.scatter_matrix(frame, alpha=0.5, figsize=None, ax=None,  
grid=False, diagonal='hist', marker='.', density_kws=None, hist_kws=None,  
range_padding=0.05, **kwargs)
```

[\[source\]](#)

Draw a matrix of scatter plots.

Parameters:

frame : DataFrame

alpha : float, optional

Amount of transparency applied.

figsize : (float,float), optional

A tuple (width, height) in inches.

ax : Matplotlib axis object, optional

grid : bool, optional

Setting this to True will show the grid.

diagonal : ('hist', 'kde')

Pick between 'kde' and 'hist' for either Kernel Density Estimation or Histogram plot in the diagonal.

marker : str, optional

Matplotlib marker type, default '.'.

density_kws : keywords

Keyword arguments to be passed to kernel density estimate plot.

hist_kws : keywords

Keyword arguments to be passed to hist function.

range_padding : float, default 0.05

Relative extension of axis range in x and y with respect to (x_max - x_min) or (y_max - y_min).

****kwargs**

Keyword arguments to be passed to scatter function.

Returns:

numpy.ndarray

A matrix of scatter plots.

```
import pandas as pd  
import numpy as np
```

```
#make this example reproducible  
np.random.seed(0)
```


```
#create DataFrame
```

```
df = pd.DataFrame({'points': np.random.randn(1000),  
                  'assists': np.random.randn(1000),  
                  'rebounds': np.random.randn(1000)})
```

```
#view first five rows of DataFrame  
df.head()
```

	points	assists	rebounds
0	1.764052	0.555963	-1.532921
1	0.400157	0.892474	-1.711970
2	0.978738	-0.422315	0.046135
3	2.240893	0.104714	-0.958374
4	1.867558	0.228053	-0.080812

```
pd.plotting.scatter_matrix(df.iloc[:, 0:2])
```





Part III : ttest_ind

```
t_val, p_val = ttest_ind(df1, df_2)
```

ttest_ind used to check whether the unknown population means of given pair of groups are equal.

tt allows one to test the null hypothesis that the means of two groups are equal

scipy.stats.ttest_ind

```
scipy.stats.ttest_ind(a, b, axis=0, equal_var=True, nan_policy='propagate',  
permutations=None, random_state=None, alternative='two-sided', trim=0, *, keepdims=False)  
\[source\]
```

Calculate the T-test for the means of *two independent* samples of scores.

This is a test for the null hypothesis that 2 independent samples have identical average (expected) values. This test assumes that the populations have identical variances by default.

Parameters: a, b : *array_like*

The arrays must have the same shape, except in the dimension corresponding to axis (the first, by default).

axis : *int or None, default: 0*

If an int, the axis of the input along which to compute the statistic. The statistic of each axis-slice (e.g. row) of the input will appear in a corresponding element of the output. If *None*, the input will be raveled before computing the statistic.

equal_var : *bool, optional*

If True (default), perform a standard independent 2 sample test that assumes equal population variances [1]. If False, perform Welch's t-test, which does not assume equal population variance [2].

i New in version 0.11.0.

nan_policy : {'propagate', 'omit', 'raise'}

Defines how to handle input NaNs.

- *propagate*: if a NaN is present in the axis slice (e.g. row) along which the statistic is computed, the corresponding entry of the output will be NaN.
- *omit*: NaNs will be omitted when performing the calculation. If insufficient data remains in the axis slice along which the statistic is computed, the corresponding entry of the output will be NaN.
- *raise*: if a NaN is present, a *ValueError* will be raised.



Part III : patsy.dmatrices

outcome_1, predictors_1 =
patsy.dmatrices('y ~ x', df_sub)

Patsy uses R-style formulas to define the model and takes care of transforming the data for you.

We can pass the relation among variables as strings.

e.g.

'y~x'

Or

'y~ x1 + x2 + b*x3'

4. Patsy

[Patsy](#) is a neat API to transform your data into experimentation model form. For regression and classification problems, you often want your data in the xy form where x is a matrix (independent variable) and y is a column vector (dependent variable). In regression, let's say you have x_1, x_2 as your independent variable and y as your dependent variable. You might want to express the possible models as follows.

- $y = b_0 + x_1 + x_2$
- $y = b_0 + x_1 + x_2 + x_1x_2$
- $y = b_0 + x_1 + x_2 + x_1^2 + x_2^2$
- $y = b_0 + x_1 + x_2 + x_1^2 + x_2^2 + x_1x_2$

```
[3]: from patsy import dmatrices

formula = 'y ~ height + weight + I(height**2) + I(weight**2) + height:weight'
y, X = dmatrices(formula, df, return_type='dataframe')
```

X

```
[3]:
```

	Intercept	height	weight	I(height ** 2)	I(weight ** 2)	height:weight
0	1.0	10.0	88.0	100.0	7744.0	880.0
1	1.0	20.0	99.0	400.0	9801.0	1980.0
2	1.0	30.0	125.0	900.0	15625.0	3750.0
3	1.0	40.0	155.0	1600.0	24025.0	6200.0
4	1.0	50.0	120.0	2500.0	14400.0	6000.0



Part III : statsmodels.regression.linear_model.OLS

```
# Now use statsmodels to initialize an OLS linear model
# This step initializes the model, and provides the data (but
# does not actually compute the model)
mod_log = sm.OLS(outcome, predictors)

# fit the model
res_log = mod_log.fit()

# Check out the results
print(res_log.summary())
```

Introduction :

A linear regression model establishes the relation between a dependent variable(**y**) and at least one independent variable(**x**) as :

$$\hat{y} = b_1x + b_0$$

In *OLS* method, we have to choose the values of b_1 and b_0 such that, the total sum of squares of the difference between the calculated and observed values of y , is minimised.

Formula for OLS:

$$S = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - b_1x_i - b_0)^2 = \sum_{i=1}^n (\epsilon_i)^2 = \min$$

Where,

\hat{y}_i = predicted value for the i th observation

y_i = actual value for the i th observation

ϵ_i = error/residual for the i th observation

n = total number of observations

To get the values of b_0 and b_1 which minimise S , we can take a partial derivative for each coefficient and equate it to zero.



THANKS!

Questions on Campuswire or office hours

Office hours: Tue/Thu, 4-5 PM

